

## Advanced Classification: Classifiers and Support Vector Machine

### Support vector classifier

The **e1071** library contains implementations for a number of statistical learning methods. In particular, the **svm()** function can be used to fit a support vector classifier when the argument **kernel="linear"** is used. A **cost** argument allows us to specify the cost of a violation to the margin. When the **cost** argument is small, then the margins will be wide and many support vectors will be on the margin or will violate the margin. When the **cost** argument is large, then the margins will be narrow and there will be few support vectors on the margin or violating the margin.

In this project we assume that students will implement independently all necessary steps like setting the working directory and connecting the library, which were explained before in HW1 and HW2.


**Step 1** Set variable **k** equal to the last 4 digits of your student number. Then initialize the random number generator as **set.seed(k)**. This is an important requirement which makes all project results different for all students with very high level of probability. Do not re-set this value for other steps of this work.

```
> set.seed(6388)
```

**Step 2** (1 mark) We begin by generating the observations, which belong to two classes, and checking whether the classes are linearly separable. Use commands **matrix** to generate two sets of data.

Plot these data using command **plot**. Demonstrate this plot and answer to the questions if these two sets are separable.

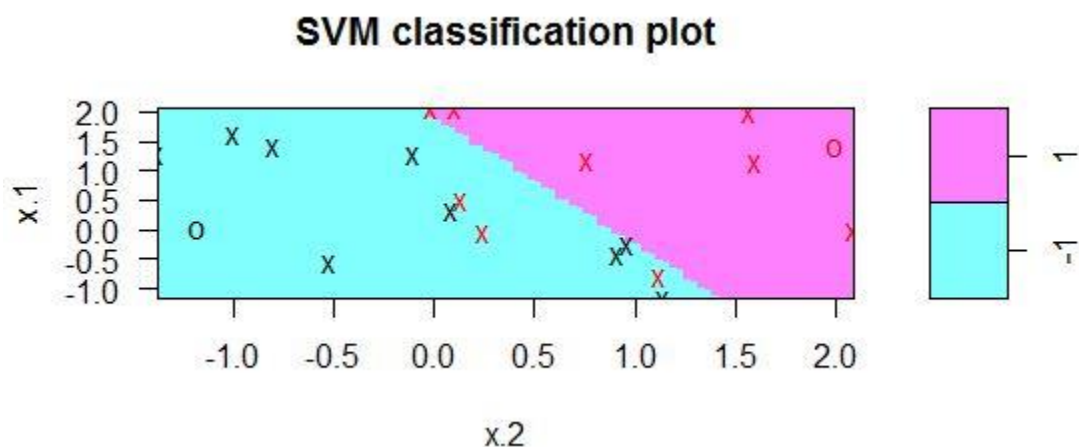
```
> x <- matrix(rnorm(20*2), ncol=2)
> y <- c(rep(-1,10), rep(1,10))
> x[y==1,]=x[y==1,] + 1
> plot(x, col=(3-y))
```



**Step 3** (1 mark) Fit the support vector classifier for cost function value 0.1. Note that in order for the `svm()` function to perform classification (as opposed to SVM-based regression), we must encode the response as a factor variable. Provide summary of the `svmfit`. Plot the support vector classifier obtained.

The important point is that before following the instructions from the text book, or use the R commands from the website, you have to install package [e1071](#).

```
> library(e1071)
> dat <- data.frame(x=x, y=as.factor(y))
> svm.fit<- svm(y ~., data=dat, kernel = 'linear', cost=0.1, scale=FALSE)
> plot(svm.fit, dat)
```



**Step 4** (1 mark) Determine their identities of the support vectors.

```
> summary(svm.fit)
```

```
Call:
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 0.1,
     scale = FALSE)
```

Parameters:  
 SVM-Type: C-classification  
 SVM-Kernel: linear  
 cost: 0.1  
 gamma: 0.5

Number of Support Vectors: 18  
 ( 9 9 )

Number of Classes: 2


Levels:  
 -1 1

```
> svm.fit$index
[1] 1 2 3 4 5 6 8 9 10 11 12 13 14 15 16 17 18 19
```

The summary lets us know there were 18 support vectors which are {1 2 3 4 5 6 8 9 10 11 12 13 14 15 16 17 18 19},  
 9 in the first class and 9 in the second

**Step 5** (1 mark) Increase number of cost parameter to 10. Check and identify the support vectors, wrote how they number changed.

```
> dat <- data.frame(x=x, y=as.factor(y))
> svm.fit1 <- svm(y ~., data=dat, kernel='linear', cost=10, scale=FALSE)
> plot(svm.fit1, dat)
```



```
> summary(svm.fit1)
```

Call:  
 svm(formula = y ~ ., data = dat, kernel = "linear",  
 cost = 10,  
 scale = FALSE)

Parameters:  
 SVM-Type: C-classification  
 SVM-Kernel: linear

```
cost: 10
gamma: 0.5
```

Number of Support Vectors: 9

```
( 5 4 )
```

Number of Classes: 2

```
Levels:
-1 1
```

```
> svm.fit1$index
[1] 1 2 3 4 9 11 16 17 19
```

```
>
```

The summary lets us know there were 18 support vectors which are {1 2 3 4 9 11 16 17 19},

5 in the first class and 4 in the second

**Step 6** (1 mark) Compare SVMs with a linear kernel, using a range of values of the **cost** parameter. Print and interpret summary.

```
> set.seed(6388)
>
> tune.out <- tune(svm, y ~., data=dat, kernel='linear',
+                 ranges=list(cost=c(0.001,0.01,0.1,1,5,10,100)))
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
cost
1
```

- best performance: 0.25

- Detailed performance results:

	cost	error	dispersion
1	1e-03	0.70	0.4216370
2	1e-02	0.70	0.4216370
3	1e-01	0.30	0.3496029
4	1e+00	0.25	0.3535534
5	5e+00	0.25	0.3535534
6	1e+01	0.25	0.3535534
7	1e+02	0.25	0.3535534

#The best cost is 1 for the output.

# best performance: 0.25

**Step 7** (1 mark) The **tune()** function stores the best model obtained; accessed it using the command. Print summary.

```
> bestmod = tune.out$best.model
```

```
> summary(bestmod)
```

```
Call:
best.tune(method = svm, train.x = y ~ ., data = dat, ranges = list(cost = c(0.001,
  0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
```

```
Parameters:
  SVM-Type:  C-classification
  SVM-Kernel: linear
    cost:    1
   gamma:   0.5
```

```
Number of Support Vectors: 13
( 7 6 )
```

```
Number of Classes: 2
```

```
Levels:
-1 1
```

Here we see that cost= 1 results in the lowest cross-validation error rate.

**Step 8** (2 marks) Generate the test data set and predict the class labels of these test observations.

```
> xtest=matrix(rnorm(20*2), ncol=2)
> ytest=sample(c(-1,1), 20, rep=TRUE)
> xtest [ ytest ==1 ,]= xtest [ ytest ==1 ,] + 1
> testdat=data.frame(x=xtest, y=as.factor(ytest))
> yhat <- predict(tune.out$best.model, testdat)
> #install.packages("caret")
> library(caret)
> confusionMatrix(yhat, testdat$y)
Confusion Matrix and Statistics
```

	Reference	
Prediction	-1	1
-1	5	2
1	3	10

Accuracy : 0.75  
95% CI : (0.509, 0.9134)  
No Information Rate : 0.6  
P-Value [Acc > NIR] : 0.1256


Kappa : 0.4681  
McNemar's Test P-Value : 1.0000

Sensitivity : 0.6250  
Specificity : 0.8333  
Pos Pred Value : 0.7143  
Neg Pred Value : 0.7692  
Prevalence : 0.4000  
Detection Rate : 0.2500  
Detection Prevalence : 0.3500  
Balanced Accuracy : 0.7292

'Positive' Class : -1

**Step 9** (2 marks) Now consider a situation in which the two classes are linearly separable. Then find a separating hyperplane using the `svm()` function. Separate the two classes in our simulated data so that they are linearly separable.

```
> x[y==1 ,]= x[y==1 ,]+0.1
> plot(x, col =(y+5) /2, pch =19)
```



**Step 10** (2 marks) Fit the support vector classifier and plot the resulting hyperplane, using a very large value of `cost` so that no observations are misclassified.

```
> dat=data.frame(x=x,y=as.factor (y))
> svmfit =svm(y~ ., data=dat , kernel ="linear", cost =1e5)
> summary (svmfit)
```

Call:  
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 1e+05)

Parameters:  
SVM-Type: C-classification  
SVM-Kernel: linear  
cost: 1e+05  
gamma: 0.5


Number of Support Vectors: 3  
( 2 1 )

Number of Classes: 2

Levels:  
-1 1

#No of supporting vectors is 3

```
> plot(svmfit , dat)
```



**Step 11** (1 marks) Answer the multiple choice question:

1. Are the support vectors outside of the margin?
2. Are the support vectors on the boarder of the margin?
3. Are the support vectors within the margin?

- #1. Are the support vectors outside of the margin? Yes  
 #2. Are the support vectors on the boarder of the margin? Yes  
 #3. Are the support vectors within the margin? No


## Support vector machine (Refer Section 9.6 from the text book) 5 marks

In order to fit an SVM using a non-linear kernel, use the `svm()` function. Use a different value of the parameter `kernel`. To fit an SVM with a polynomial kernel use `kernel="polynomial"`, and to fit an SVM with a radial kernel use `kernel="radial"`. In the former case we also use the `degree` argument to specify a degree for the polynomial kernel (this is  $d$  in (9.22)), and in the latter case we use `gamma` to specify a value of  $\gamma$  for the radial basis kernel (9.24).

**Step 1** (1 marks) Generate some data with a non-linear class boundary and plot them.

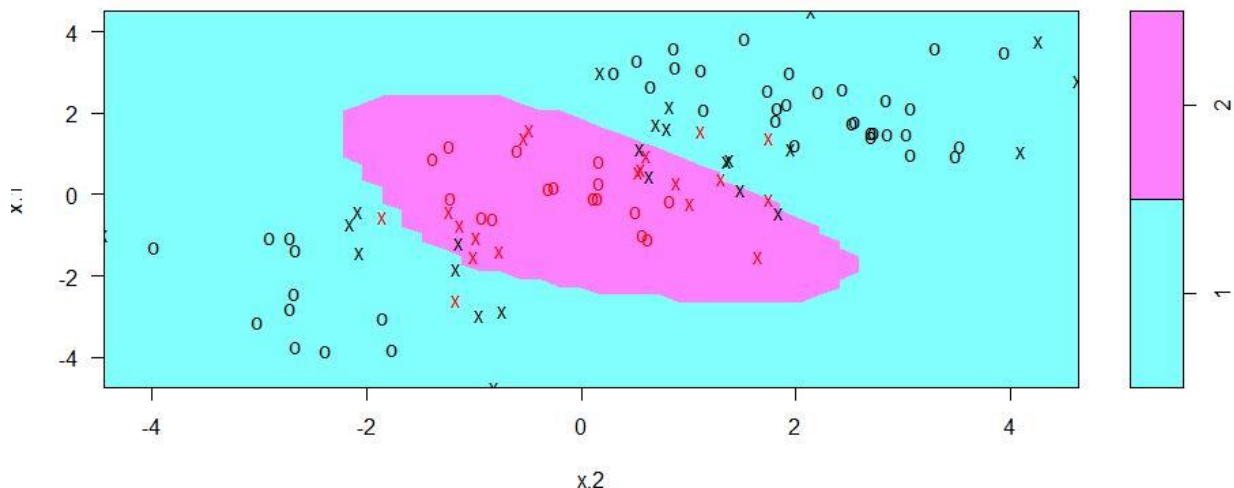
```
> plot(svmfit , dat)
> set.seed (6388)
> x=matrix (rnorm (200*2) , ncol =2)
> x[1:100 ,]=x[1:100 ,]+2
> x[101:150 ,]= x[101:150 ,] -2
> y=c(rep (1 ,150) ,rep (2 ,50) )
```

```
> dat=data.frame(x=x,y=as.factor (y))
> plot(x, col=y)
```



**Step 2** (1 marks) Fit the training data using the svm() function with a radial kernel and  $\gamma = 1$ .

```
> train=sample (200 ,100)
> svmfit =svm(y~., data=dat [train, ], kernel = "radial", gamma =1,cost =1)
> plot(svmfit , dat[train,])
```



**Step 3** (1 marks) Print summary. What can you tell about of the error? Re-fit the SVM classification with higher cost. Print summary and plot results. What are your major concern about these results?

```
> summary(svmfit)
```

```
Call:
svm(formula = y ~ ., data = dat[train, ], kernel = "radial",
     gamma = 1, cost = 1)
```

Parameters:  
SVM-Type: c-classification



```
SVM-kernel:  radial
           cost:  1
           gamma: 1
```

Number of Support Vectors: 43


( 19 24 )

Number of Classes: 2

```
Levels:
 1 2
```

#No of support vectors is 43.

```
> svmfit =svm(y~., data=dat [train ,], kernel ="radial",gamma =1, cost=1e5)
> plot(svmfit ,dat [train ,])
```



#We can see from the figure that there are a fair number of training errors in this SVM fit. If we increase the value of cost, we can reduce the number of training errors. However, this comes at the price of a more irregular decision boundary that seems to be at risk of overfitting the data.

**Step 4** (1 marks) Perform cross-validation using `tune()` to select the best choice of  $\gamma$  and cost for an SVM with a radial kernel.

```
> set.seed (6388)
> tune.out=tune(svm , y~., data=dat[train ,], kernel
="radial", ranges =list(cost=c(0.1 ,1 ,10 ,100 ,1000),
gamma=c(0.5,1,2,3,4) ))
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation
- best parameters:  
cost gamma

1 0.5

- best performance: 0.1

- Detailed performance results:

	cost	gamma	error	dispersion
1	1e-01	0.5	0.15	0.08498366
2	1e+00	0.5	0.10	0.04714045
3	1e+01	0.5	0.12	0.06324555
4	1e+02	0.5	0.11	0.05676462
5	1e+03	0.5	0.16	0.08432740
6	1e-01	1.0	0.16	0.09660918
7	1e+00	1.0	0.11	0.05676462
8	1e+01	1.0	0.12	0.06324555
9	1e+02	1.0	0.16	0.08432740
10	1e+03	1.0	0.21	0.11972190
11	1e-01	2.0	0.28	0.11352924
12	1e+00	2.0	0.13	0.08232726
13	1e+01	2.0	0.13	0.06749486
14	1e+02	2.0	0.19	0.09944289
15	1e+03	2.0	0.19	0.11005049
16	1e-01	3.0	0.29	0.11972190
17	1e+00	3.0	0.13	0.08232726
18	1e+01	3.0	0.14	0.06992059
19	1e+02	3.0	0.17	0.08232726
20	1e+03	3.0	0.22	0.09189366
21	1e-01	4.0	0.35	0.14337209
22	1e+00	4.0	0.13	0.08232726
23	1e+01	4.0	0.15	0.08498366
24	1e+02	4.0	0.17	0.08232726
25	1e+03	4.0	0.23	0.08232726

>

Therefore, the best choice of parameters involves cost=1 and gamma=0.5

**Step 5** (1 marks) Interpret results: what si the optimal values of cost and  $\gamma$  and what is the lowestt percent of misclassified objects?

```
> yhat <- predict(tune.out$best.model, dat[-train,])
> confusionMatrix(yhat, dat[-train, 'y'])
Confusion Matrix and Statistics
```

	Reference	
Prediction	1	2
1	75	0
2	10	15

Accuracy : 0.9  
95% CI : (0.8238, 0.951)  
No Information Rate : 0.85  
P-Value [Acc > NIR] : 0.099447

Kappa : 0.6923  
McNemar's Test P-Value : 0.004427

Sensitivity : 0.8824  
Specificity : 1.0000  
Pos Pred Value : 1.0000  
Neg Pred Value : 0.6000

```
Prevalence : 0.8500
Detection Rate : 0.7500
Detection Prevalence : 0.7500
Balanced Accuracy : 0.9412
```

```
'Positive' Class : 1
```

>

The optimal values of cost is 1 and gamma=0.5. Percentage of misclassified objects is 10 percent.

## Decision trees for classification (Refer Section 8.3 from the text book)

7 marks

**Step 1** The **ISLR** and **tree** libraries are used to construct classification and regression trees. First use classification trees to analyze the **Carseats** data set. In these data, **Sales** is a continuous variable, and so we begin by recoding it as a binary variable. Use the **ifelse()** function to create a variable, called **High**, which takes on a value of **Yes** if the **Sales** variable exceeds 8, and takes on a value of **No** otherwise. Do not forget to install relevant packages. The description of ISLR package including **Carseats** (which contains **Sales**) data set is available on the course website (R language page).

```
> #install.packages('tree', dependencies=TRUE)
> library (tree)
warning message:
package 'tree' was built under R version 3.4.3
> library (ISLR)
```

>

**Step 2** (1 marks) Use the **data.frame()** function to merge **High** with the rest of the **Carseats** data. Use the **tree()** function to fit a classification tree in order to predict **High** using all variables but **Sales**. The syntax of the **tree()** function is quite similar to that of the **lm()** function. Use **summary()** function lists the variables that are used as internal nodes in the tree, the number of terminal nodes, and the (training) error rate. What is the training error rate?

```
> attach (Carseats )
The following objects are masked from Carseats (pos =
3):
```

```
Advertising, Age, CompPrice, Education, Income,
Population, Price, Sales, ShelfLoc, Urban, US
```

```
> View(Carseats)
> High=ifelse (Sales <=8," No"," Yes ")
> Carseats =data.frame(Carseats ,High)
> tree.carseats =tree(High~.-Sales ,Carseats )
> summary (tree.carseats )
```

Classification tree:

```
tree(formula = High ~ . - Sales, data = Carseats)
Variables actually used in tree construction:
```

```

[1] "ShelveLoc" "Price" "Income"
"CompPrice"
[5] "Population" "Advertising" "Age" "US"
Number of terminal nodes: 27
Residual mean deviance: 0.4575 = 170.7 / 373
Misclassification error rate: 0.09 = 36 / 400

```


>

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban	US	High	High.1
1	9.50	138	73	11	276	120	Bad	42	17	Yes	Yes	Yes	Yes
2	11.22	111	48	16	260	83	Good	65	10	Yes	Yes	Yes	Yes
3	10.06	113	35	10	269	80	Medium	59	12	Yes	Yes	Yes	Yes
4	7.40	117	100	4	466	97	Medium	55	14	Yes	Yes	No	No
5	4.15	141	64	3	340	128	Bad	38	13	Yes	No	No	No
6	10.81	124	113	13	501	72	Bad	78	16	No	Yes	Yes	Yes
7	6.63	115	105	0	45	108	Medium	71	15	Yes	No	No	No
8	11.85	136	81	15	425	120	Good	67	10	Yes	Yes	Yes	Yes
9	6.54	132	110	0	108	124	Medium	76	10	No	No	No	No
10	4.69	132	113	0	131	124	Medium	76	17	No	Yes	No	No
11	9.01	121	78	9	150	100	Bad	26	10	No	Yes	Yes	Yes
12	11.96	117	94	4	503	94	Good	50	13	Yes	Yes	Yes	Yes

Mis classification rate is .09

**Step 3** (1 marks) Plot and text the car seat tree. Provide in your answer the tree **without** texts.

> plot(tree.carseats )



**Step 4** (1 marks) Type the name of the tree object, and analyze the **R** prints output corresponding to each branch of the tree. **R** displays the split criterion (e.g. **Price**<92.5), the number of observations in that branch, the deviance, the overall prediction for the branch (**Yes** or **No**), and the fraction of observations in that branch that take on values of **Yes** and **No**. Branches that lead to terminal nodes are indicated using asterisks.

> **tree.carseats**

node), split, n, deviance, yval, (yprob)

\* denotes terminal node

```

1) root 400 541.500 No ( 0.59000 0.41000 )
2) ShelfLoc: Bad,Medium 315 390.600 No ( 0.68889 0.31111 )
4) Price < 92.5 46 56.530 Yes ( 0.30435 0.69565 )
8) Income < 57 10 12.220 No ( 0.70000 0.30000 )
16) CompPrice < 110.5 5 0.000 No ( 1.00000 0.00000 ) *
17) CompPrice > 110.5 5 6.730 Yes ( 0.40000 0.60000 ) *
9) Income > 57 36 35.470 Yes ( 0.19444 0.80556 )
18) Population < 207.5 16 21.170 Yes ( 0.37500 0.62500 ) *
19) Population > 207.5 20 7.941 Yes ( 0.05000 0.95000 ) *
5) Price > 92.5 269 299.800 No ( 0.75465 0.24535 )
10) Advertising < 13.5 224 213.200 No ( 0.81696 0.18304 )
20) CompPrice < 124.5 96 44.890 No ( 0.93750 0.06250 )
40) Price < 106.5 38 33.150 No ( 0.84211 0.15789 )
80) Population < 177 12 16.300 No ( 0.58333 0.41667 )
160) Income < 60.5 6 0.000 No ( 1.00000 0.00000 ) *
161) Income > 60.5 6 5.407 Yes ( 0.16667 0.83333 ) *
81) Population > 177 26 8.477 No ( 0.96154 0.03846 ) *
41) Price > 106.5 58 0.000 No ( 1.00000 0.00000 ) *
21) CompPrice > 124.5 128 150.200 No ( 0.72656 0.27344 )
42) Price < 122.5 51 70.680 Yes ( 0.49020 0.50980 )
84) ShelfLoc: Bad 11 6.702 No ( 0.90909 0.09091 ) *
85) ShelfLoc: Medium 40 52.930 Yes ( 0.37500 0.62500 )
170) Price < 109.5 16 7.481 Yes ( 0.06250 0.93750 ) *
171) Price > 109.5 24 32.600 No ( 0.58333 0.41667 )
342) Age < 49.5 13 16.050 Yes ( 0.30769 0.69231 ) *
343) Age > 49.5 11 6.702 No ( 0.90909 0.09091 ) *
```

```

43) Price > 122.5 77 55.540 No ( 0.88312 0.11688 )
86) CompPrice < 147.5 58 17.400 No ( 0.96552 0.03448 ) *
87) CompPrice > 147.5 19 25.010 No ( 0.63158 0.36842 )
174) Price < 147 12 16.300 Yes ( 0.41667 0.58333 )
348) CompPrice < 152.5 7 5.742 Yes ( 0.14286 0.85714 ) *
349) CompPrice > 152.5 5 5.004 No ( 0.80000 0.20000 ) *
175) Price > 147 7 0.000 No ( 1.00000 0.00000 ) *
11) Advertising > 13.5 45 61.830 Yes ( 0.44444 0.55556 )
22) Age < 54.5 25 25.020 Yes ( 0.20000 0.80000 )
44) CompPrice < 130.5 14 18.250 Yes ( 0.35714 0.64286 )
88) Income < 100 9 12.370 No ( 0.55556 0.44444 ) *
89) Income > 100 5 0.000 Yes ( 0.00000 1.00000 ) *
45) CompPrice > 130.5 11 0.000 Yes ( 0.00000 1.00000 ) *
23) Age > 54.5 20 22.490 No ( 0.75000 0.25000 )
46) CompPrice < 122.5 10 0.000 No ( 1.00000 0.00000 ) *
47) CompPrice > 122.5 10 13.860 No ( 0.50000 0.50000 )
94) Price < 125 5 0.000 Yes ( 0.00000 1.00000 ) *
95) Price > 125 5 0.000 No ( 1.00000 0.00000 ) *
3) ShelveLoc: Good 85 90.330 Yes ( 0.22353 0.77647 )
6) Price < 135 68 49.260 Yes ( 0.11765 0.88235 )
12) US: No 17 22.070 Yes ( 0.35294 0.64706 )
24) Price < 109 8 0.000 Yes ( 0.00000 1.00000 ) *
25) Price > 109 9 11.460 No ( 0.66667 0.33333 ) *
13) US: Yes 51 16.880 Yes ( 0.03922 0.96078 ) *
7) Price > 135 17 22.070 No ( 0.64706 0.35294 )
14) Income < 46 6 0.000 No ( 1.00000 0.00000 ) *
15) Income > 46 11 15.160 Yes ( 0.45455 0.54545 ) *

```

**Step 5** (1 marks) Evaluate the performance of a classification tree on these data and the training error. Split the observations into a training set (200 records) and a test set, build the tree using the training set, and evaluate its performance on the test data. The `predict()` function can be used for this purpose. In the case of a classification tree, the argument `type="class"` instructs `R` to return the actual class prediction.

```

> set.seed (6388)
> train=sample (1: nrow(Carseats), 200)
> Carseats.test=Carseats [-train ,]
> High.test=High[-train ]
> tree.carseats =tree(High~.-Sales ,Carseats ,subset =train )
> tree.pred=predict (tree.carseats ,Carseats.test ,type ="class")
> table(tree.pred ,High.test)

```

	High.test	
tree.pred	No	Yes
No	95	26
Yes	26	53

#This approach leads to correct predictions for around 74% of the locations in the test data set.  
 #(26+26) /200 = 0.26 , 26% is the error date.

**Step 6** (1 marks) Consider whether pruning the tree might lead to improved results. The function `cv.tree()` performs cross-validation in order to `cv.tree()` determine the optimal level of

tree complexity; cost complexity pruning is used in order to select a sequence of trees for consideration. Use the argument `FUN=prune.misclass` in order to indicate that we want the classification error rate to guide the cross-validation and pruning process, rather than the default for the `cv.tree()` function, which is deviance. The `cv.tree()` function reports the number of terminal nodes of each tree considered (`size`) as well as the corresponding error rate and the value of the cost-complexity parameter used (`k`, which corresponds to  $\alpha$  in (8.4)).

What is the optimal pruning (optimal number of leaves)?

```
> set.seed(6388)
> cv.carseats = cv.tree(tree.carseats ,FUN=prune.misclass )
> names(cv.carseats)
[1] "size"    "dev"     "k"       "method"
> cv.carseats
$size
[1] 17 13 11  9  8  4  1

$dev
[1] 66 65 65 62 56 56 87

$k
[1] -Inf  0.00  0.50  2.50  3.00  3.25 14.00



$method
[1] "misclass"

attr("class")
[1] "prune"          "tree.sequence"
```

Note that, despite the name, `dev` corresponds to the cross-validation error rate in this instance. The tree with 4 terminal nodes results in the lowest cross-validation error rate, with 56 cross-validation errors.

**Step 7** (1 marks) Plot the error rate as a function of both `size` and `k`.

```
> par(mfrow =c(1,2))
> plot(cv.carseats$size ,cv.carseats$dev ,type="b")
> plot(cv.carseats$k ,cv.carseats$dev ,type="b")
>
```



**Step 8** (1 marks) Apply the `prune.misclass()` function in order to prune the tree to **prune**. Obtain the nine-node tree. Plot it **with** text (do not care about overlapping!).

```
> prune.carseats =prune.misclass(tree.carseats,best =9)
> plot(prune.carseats)
> text(prune.carseats,pretty=0)
```



ShelveLoc: Bad

